

# Homework 1

1. The Iowa data set `iowa.csv` is a toy example that summarises the yield of wheat (bushels per acre) for the state of Iowa between 1930-1962. In addition to yield, year, rainfall and temperature were recorded as the main predictors of yield.
  - a. First, we need to load the data set into R using the command `read.csv()`. Use the help function to learn what arguments this function takes. Once you have the necessary input, load the data set into R and make it a data frame called `iowa.df`.
  - b. How many rows and columns does `iowa.df` have?
  - c. What are the names of the columns of `iowa.df`?
  - d. What is the value of row 5, column 7 of `iowa.df`?
  - e. Display the second row of `iowa.df` in its entirety.

```
iowa.df<-read.csv("data/iowa.csv",header=T)
iowa.df<-separate(iowa.df,Year.Rain0.Temp1.Rain1.Temp2.Rain2.Temp3.Rain3.Temp4.Yield,
into = c('Year','Rain0','Temp1','Rain1','Temp2','Rain2','Temp3','Rain3','Temp4','Yield'),
sep = ";")
dim(iowa.df)#answer of b
```

```
## [1] 33 10
```

```
colnames(iowa.df)#answer of c
```

```
## [1] "Year" "Rain0" "Temp1" "Rain1" "Temp2" "Rain2" "Temp3" "Rain3" "Temp4"
## [10] "Yield"
```

```
iowa.df[5,7]#answer of d
```

```
## [1] "79.7"
```

```
iowa.df[2,]#answer of e
```

```
##   Year Rain0 Temp1 Rain1 Temp2 Rain2 Temp3 Rain3 Temp4 Yield
## 2 1931 14.76  57.5   3.83    75   2.72  77.2   3.3  72.6  32.9
```

2. Syntax and class-typing.

- a. For each of the following commands, either explain why they should be errors, or explain the non-erroneous result.

```
vector1 <- c("5", "12", "7", "32")
max(vector1)
[1] "7"
non-erroneous result because vector1 is a char vec.
```

```
sort(vector1)
[1] "12" "32" "5"  "7"
non-erroneous result because vector1 is a char vec.
sum(vector1)
Error in sum(vector1): 'type'(character) Wrong parameter
error because vector1 is a char vec.
```

b. For the next series of commands, either explain their results, or why they should produce errors.

```
vector2 <- c("5",7,12)
vector2[2] + vector2[3]
Error in vector2[2] + vector2[3] :
Non numeric parameter in binary column operator
because vector2 is a char vec.
```

```
dataframe3 <- data.frame(z1="5",z2=7,z3=12)
dataframe3[1,2] + dataframe3[1,3]
[1] 19
because it is a dataframe.
```

```
list4 <- list(z1="6", z2=42, z3="49", z4=126)
list4[[2]]+list4[[4]]
[1] 168
because it is a list and [[ ]] drops names and structures.
list4[2]+list4[4]
Error in list4[2] + list4[4] :
Non numeric parameter in binary column operator
because it is a list and [] does not drop names and structures.
```

### 3. Working with functions and operators.

- The colon operator will create a sequence of integers in order. It is a special case of the function `seq()` which you saw earlier in this assignment. Using the help command `?seq` to learn about the function, design an expression that will give you the sequence of numbers from 1 to 10000 in increments of 372. Design another that will give you a sequence between 1 and 10000 that is exactly 50 numbers in length.
- The function `rep()` repeats a vector some number of times. Explain the difference between `rep(1:3, times=3)` and `rep(1:3, each=3)`.

```
seq(1,10000,372)#answer of a
```

```
## [1]      1   373   745 1117 1489 1861 2233 2605 2977 3349 3721 4093 4465 4837 5209
## [16] 5581 5953 6325 6697 7069 7441 7813 8185 8557 8929 9301 9673
```

```
seq(1,10000,length.out = 50)#answer of a
```

```
## [1]      1.0000    205.0612    409.1224    613.1837    817.2449   1021.3061
## [7]  1225.3673   1429.4286   1633.4898   1837.5510   2041.6122   2245.6735
## [13]  2449.7347   2653.7959   2857.8571   3061.9184   3265.9796   3470.0408
## [19]  3674.1020   3878.1633   4082.2245   4286.2857   4490.3469   4694.4082
## [25]  4898.4694   5102.5306   5306.5918   5510.6531   5714.7143   5918.7755
## [31]  6122.8367   6326.8980   6530.9592   6735.0204   6939.0816   7143.1429
```

```
## [37] 7347.2041 7551.2653 7755.3265 7959.3878 8163.4490 8367.5102
## [43] 8571.5714 8775.6327 8979.6939 9183.7551 9387.8163 9591.8776
## [49] 9795.9388 10000.0000
```

```
rep(1:3,times=3)#answer of b
```

```
## [1] 1 2 3 1 2 3 1 2 3
```

```
rep(1:3,each=3)#answer of b
```

```
## [1] 1 1 1 2 2 2 3 3 3
```

Explain:

`rep(1:3,times=3)` means `seq(1,3)` is repeated 3 times.

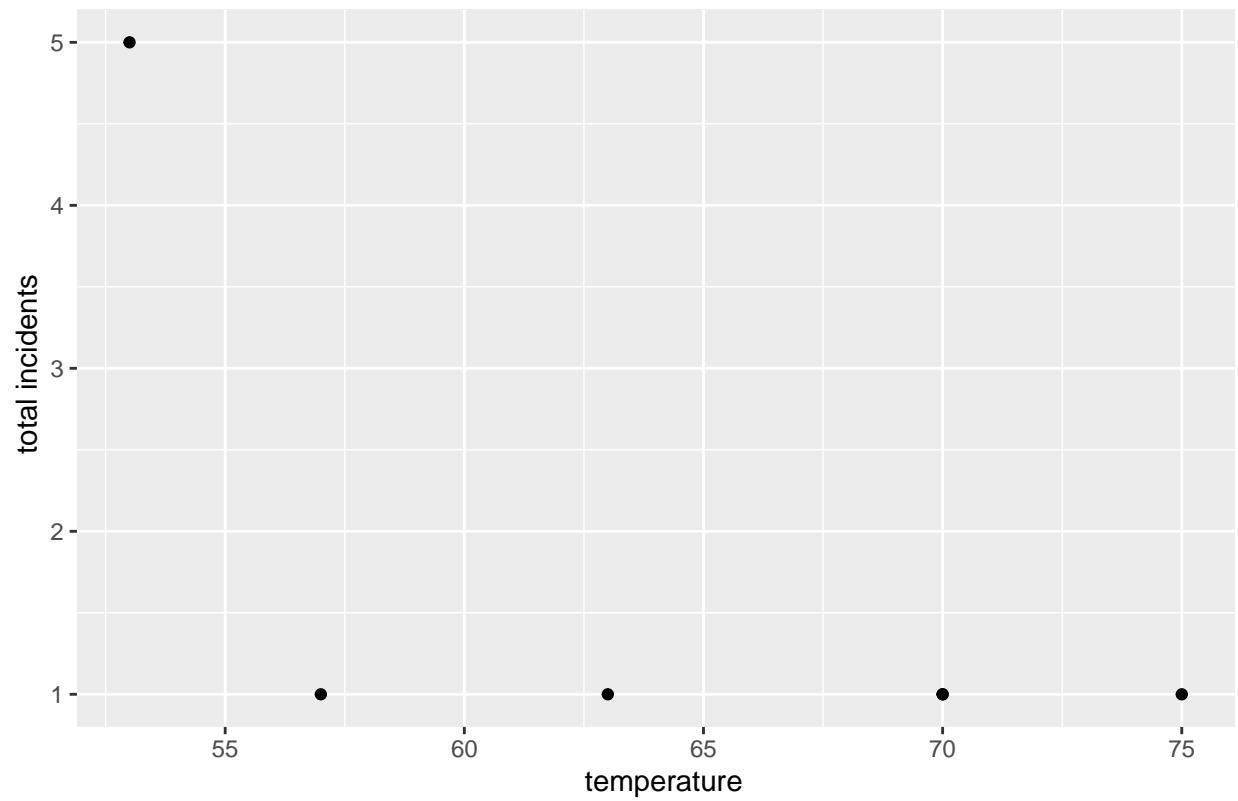
`rep(1:3,each=3)` means each element of `seq(1,3)` is repeated 3 times.

MB.Ch1.2. The `orings` data frame gives data on the damage that had occurred in US space shuttle launches prior to the disastrous Challenger launch of 28 January 1986. The observations in rows 1, 2, 4, 11, 13, and 18 were included in the pre-launch charts used in deciding whether to proceed with the launch, while remaining rows were omitted.

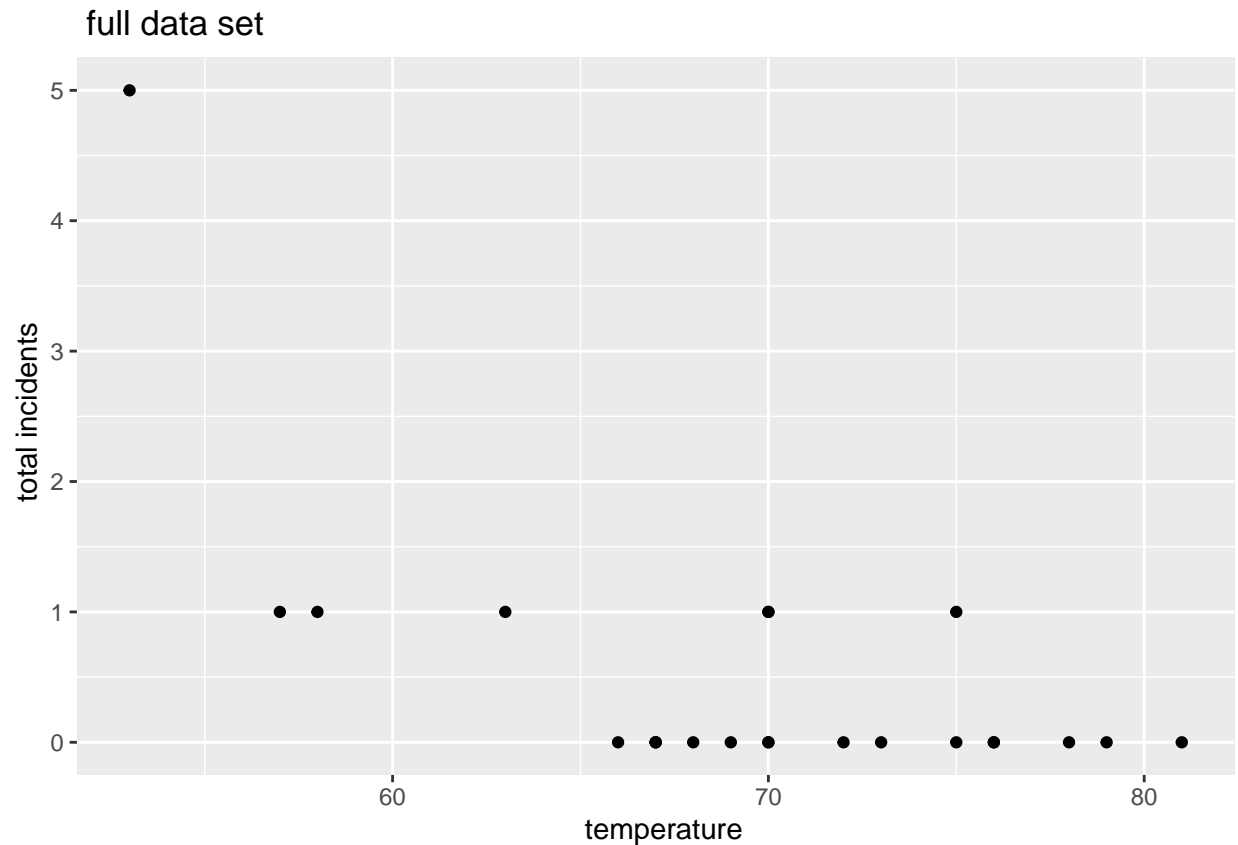
Create a new data frame by extracting these rows from `orings`, and plot total incidents against temperature for this new data frame. Obtain a similar plot for the full data set.

```
deciding<-orings[c(1,2,4,11,13,18),]
ggplot(data = deciding) +
  geom_point(aes(x = Temperature, y = Total))+
  labs(x = " temperature",
       y = " total incidents",
       title = " extracting data set")
```

### extracting data set



```
ggplot(data = orings) +  
  geom_point(aes(x = Temperature, y = Total)) +  
  labs(x = " temperature",  
       y = " total incidents",  
       title = " full data set")
```



MB.Ch1.4. For the data frame ais (DAAG package)

- (a) Use the function `str()` to get information on each of the columns. Determine whether any of the columns hold missing values.

```
str(ais) # none of the columns hold missing values.
```

```
## 'data.frame':    202 obs. of  13 variables:
## $ rcc      : num  3.96 4.41 4.14 4.11 4.45 4.1 4.31 4.42 4.3 4.51 ...
## $ wcc      : num  7.5 8.3 5 5.3 6.8 4.4 5.3 5.7 8.9 4.4 ...
## $ hc       : num  37.5 38.2 36.4 37.3 41.5 37.4 39.6 39.9 41.1 41.6 ...
## $ hg       : num  12.3 12.7 11.6 12.6 14 12.5 12.8 13.2 13.5 12.7 ...
## $ ferr     : num  60 68 21 69 29 42 73 44 41 44 ...
## $ bmi      : num  20.6 20.7 21.9 21.9 19 ...
## $ ssf      : num  109.1 102.8 104.6 126.4 80.3 ...
## $ pcBfat   : num  19.8 21.3 19.9 23.7 17.6 ...
## $ lbm      : num  63.3 58.5 55.4 57.2 53.2 ...
## $ ht       : num  196 190 178 185 185 ...
## $ wt       : num  78.9 74.4 69.1 74.9 64.6 63.7 75.2 62.3 66.5 62.9 ...
## $ sex      : Factor w/ 2 levels "f","m": 1 1 1 1 1 1 1 1 1 1 ...
## $ sport    : Factor w/ 10 levels "B_Ball","Field",...: 1 1 1 1 1 1 1 1 1 1 ...
```

- (b) Make a table that shows the numbers of males and females for each different sport. In which sports is there a large imbalance (e.g., by a factor of more than 2:1) in the numbers of the two sexes?

```
ais%>%dplyr::mutate(num=1)%>%
  dplyr::group_by(sport,sex) %>%
  dplyr::summarize(cont_num = sum(num))%>%
  dplyr::ungroup()
```

```
## 'summarise()' regrouping output by 'sport' (override with '.groups' argument)
```

```
## # A tibble: 17 x 3
##   sport    sex  cont_num
##   <fct>   <fct>   <dbl>
## 1 B_Ball  f         13
## 2 B_Ball  m         12
## 3 Field   f          7
## 4 Field   m         12
## 5 Gym     f          4
## 6 Netball f         23
## 7 Row     f         22
## 8 Row     m         15
## 9 Swim    f          9
## 10 Swim   m         13
## 11 T_400m f         11
## 12 T_400m m         18
## 13 T_Sprnt f          4
## 14 T_Sprnt m         11
## 15 Tennis f          7
## 16 Tennis m          4
## 17 W_Polo m         17
```

Explain: there exists large imbalance in the numbers of the two sexes in Gym,Netball,T\_Sprnt and W\_Polo.


MB.Ch1.6.Create a data frame called Manitoba.lakes that contains the lake's elevation (in meters above sea level) and area (in square kilometers) as listed below. Assign the names of the lakes using the row.names() function. elevation area Winnipeg 217 24387 Winnipegosis 254 5374 Manitoba 248 4624 SouthernIndian 254 2247 Cedar 253 1353 Island 227 1223 Gods 178 1151 Cross 207 755 Playgreen 217 657

```
Manitoba.lakes <- data.frame(elevation=c( 217,254,248,254,253,227,178,207,217),
                             area=c(24387,5374,4624,2247,1353,1223,1151,755,657))
rownames(Manitoba.lakes) <- c("Winnipeg","Winnipegosis","Manitoba ",
  "SouthernIndian","Cedar","Island",
  "Gods","Cross","Playgreen")
Manitoba.lakes
```

```
##           elevation  area
## Winnipeg           217 24387
## Winnipegosis       254  5374
## Manitoba           248  4624
## SouthernIndian     254  2247
## Cedar              253  1353
## Island             227  1223
## Gods               178  1151
## Cross              207   755
## Playgreen          217   657
```

- (a) Use the following code to plot  $\log_2(\text{area})$  versus elevation, adding labeling information (there is an extreme value of area that makes a logarithmic scale pretty much essential):


```
attach(Manitoba.lakes)
plot(log2(area) ~ elevation, pch=16, xlim=c(170,280))
# NB: Doubling the area increases log2(area) by 1.0
text(log2(area) ~ elevation, labels=row.names(Manitoba.lakes), pos=4)
text(log2(area) ~ elevation, labels=area, pos=2)
title("Manitoba's Largest Lakes")
```



Devise captions that explain the labeling on the points and on the y-axis. It will be necessary to explain how distances on the scale relate to changes in area.

- (b) Repeat the plot and associated labeling, now plotting area versus elevation, but specifying  $\log="y"$  in order to obtain a logarithmic y-scale.

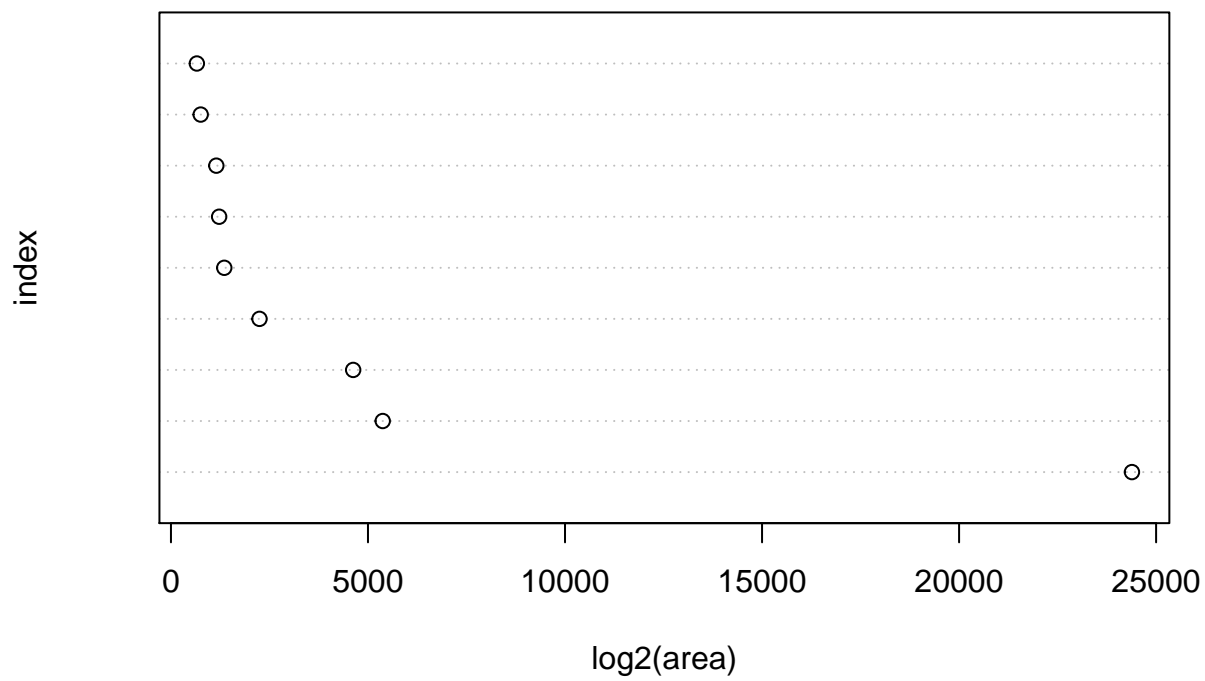
```
plot(area ~ elevation, pch=16, xlim=c(170,280), ylog=T)
text(area ~ elevation, labels=row.names(Manitoba.lakes), pos=4, ylog=T)
text(area ~ elevation, labels=area, pos=2, ylog=T)
title("Manitoba's Largest Lakes")
```




MB.Ch1.7. Look up the help page for the R function `dotchart()`. Use this function to display the areas of the Manitoba lakes (a) on a linear scale, and (b) on a logarithmic scale. Add, in each case, suitable labeling information.

```
dotchart(area, xlab = "log2(area)", ylab = "index")
```





```
dotchart(log2(area), xlab = "log2(area)", ylab = "index")
```



MB.Ch1.8. Using the `sum()` function, obtain a lower bound for the area of Manitoba covered by water.

```
sum(area)
```

```
## [1] 41771
```

i.e. a lower bound for the area of Manitoba covered by water is 41771.